Final Project

Data 102: Data, Inference, and Decisions, Spring 2021

Due Date: Monday May 10, 2021

In this project, you will complete a guided analysis for a dataset of your choice. We have curated a list of suggested datasets, but you are welcome to select an external dataset.

Your analysis should include the following steps:

1. **Data Overview** Introduce your dataset and describe the process under which it was generated.
2. **EDA** Perform exploratory data analysis (EDA) and describe key features of your dataset. You may want to complete this step before deciding on a research question below.
3. **Research Questions** List **two** research questions that you will explore in this project. Between the two research questions, you should use at least two of the four techniques that you've learned this semester. You may use more than two techniques, and you may use more than one technique for any particular question.
   - Testing many hypotheses with appropriate corrections,
   - Bayesian inference
   - Causal inference,
   - Comparing generalized linear models (GLMs) to nonparametric methods.

   At least one of your techniques should be either **Bayesian inference** or **causal inference**. Please see Section 1 for examples and clarification.
4. **Inference and Decisions** Apply the two techniques you chose above to answer your research questions, explaining your choices.
5. **Conclusion** Highlight key findings. Identify potential next steps. Assess the strengths and limitations of your analysis.

You must work in groups of four. If you don't have a group you can fill out this form by Wednesday 4/21 at 11:59PM PT and you will randomly be assigned a group. If you do have a group, **you must fill out the same form** by the deadline to declare your group. In very special circumstances, we will allow students to work alone. These exceptions are reserved for students in extraneous personal circumstances or who would like to apply class methods to an ongoing personal project (as in the case of a publication or senior thesis). If you believe you qualify for this exception, please email data102@berkeley.edu by Wednesday, 4/21 at 11:59 PM PT with relevant information/documentation. Do **not** assume the exception has been granted until you receive a confirmation email.

Detailed guidelines are provided below. Please read through this entire document before you begin working!

Contents

1 Research Question Examples **4**

2 Section Guidelines **4**
- 2.1 Data Overview . . . 4
- 2.2 Research Questions . . . 5
- 2.3 EDA . . . 5
- 2.4 Option A: Multiple hypothesis testing / decision making . . . 5
- 2.5 Option B: Bayesian Inference . . . 6
- 2.6 Option C: Prediction with GLMs and nonparametric methods . . . 6
- 2.7 Option D: Causal Inference . . . 7
- 2.8 Conclusions . . . 7

3 Project Deliverables **8**
- 3.1 Project Proposal (Due 4/30) . . . 8
- 3.2 Written Report . . . 8
- 3.3 Jupyter Notebook . . . 8
- 3.4 Dataset (for external datasets) . . . 8
- 3.5 Team Member Assessment . . . 8

4 Datasets **9**
- 4.1 Dataset 1: Chronic Disease and Air Quality . . . 9
- 4.2 Dataset 2: Transportation, Mobility, and Infrastructure . . . 10
- 4.3 Dataset 3: Primary Election Endorsements and Financing . . . 10
- 4.4 Guidelines for External Data . . . 11
- 4.5 External Data Sources . . . 11

5 Grading **12**

1 Research Question Examples

Here are some examples of research questions on hypothetical datasets. Note that all of them use either Bayesian inference or causal inference, all of them use at least two of the techniques described, and all of them answer at least two research questions.

- If you were looking at a dataset of Data 102 students, you might choose as your research questions (1) does attending office hours cause an improvement in homework grades (causal inference), and (2) Can we fit a Bayesian Gaussian mixture model to the distributions of assignment grades by student year (Bayesian inference).
- If you were looking at a dataset involving jellybean consumption, acne, and other demographics, you might choose as your research questions (1) does consuming different colors of jellybeans cause acne (causal inference *and* multiple hypothesis testing), and (2) predicting jelly bean consumption from personal demographics, using negative binomial regression and random forests (prediction with GLMs and nonparametric methods).
- If you were looking at a dataset involving characters in a TV show (lines of dialogue, gender, age, etc.), you might choose as your research questions (1) how well does character demographic information predict lines of dialogue for each season, using a Bayesian GLM and nonparametric methods (Bayesian inference *and* prediction with GLMs and nonparametrics); and (2) for each season, is there an association between gender and lines spoken (multiple hypothesis testing).

2 Section Guidelines

Your report should include each of the following sections, and address the listed questions at minimum. You should include additional, relevant discussion to each section that is specific to the features of your dataset.

Depending on your research questions, you should choose at least two of the corresponding sections for options A through D.

2.1 Data Overview

- How was your data generated? Is it a sample or census?
- If you chose to use your own data, describe the data source and download process.
- If your data represents a sample:
  - Compare the distribution of one of your variables to what is expected in the population. For example, if your data has an age variable, compare it to the age structure of the population.
    - Do you notice any differences?
    - How does this affect the generalizability of your results?
- If your data represents a census:
  - Are there any groups that were systematically excluded from your data?
- To what extent were participants aware of the collection/use of this data?

- What is the granularity of your data? What does each row represent? How will that impact the interpretation of your findings?
- Are any of the following concerns relevant in the context of your data?
  - Selection bias
  - Measurement error
  - Convenience sampling
- Are there important features/columns that you wish you had, but are unavailable? What are they and what questions would they help you answer?

2.2 Research Questions

- Your research questions should involve using the methods mentioned above (i.e., the ones you learned in Data 102) to answer them. For each research question, describe:
  - What is your first research question? What real-world decision(s) could be made by answering it?
  - Explain why the method you will use is a good fit for the question (for example, if you choose causal inference, you should explain why your research question is a good fit for causal inference).

2.3 EDA

- Visualize at least two quantitative variables and two categorical variables. Describe any trends you observe, and any relationships you may want to follow up on.
- Describe any data cleaning steps you took. How will these decisions impact your model and inferences?
  - For example, are there particular groups that have higher frequencies of missing values?

2.4 Option A: Multiple hypothesis testing / decision making

- **Methods**
  - Describe the hypotheses that you'll be testing using your dataset, and explain why it makes sense to test many hypotheses instead of just one.
  - Describe how you'll be testing each hypothesis.
  - Describe at least two different ways you'll correct for multiple hypothesis tests.
- **Results**
  - Summarize and interpret the results from the hypothesis tests themselves.
  - For the two correction methods you chose, clearly explain what kind of error/rate is being controlled by each one.
- **Discussion**
  - Which hypotheses "survived" the correction procedures? If none did, explain why.

- What decisions can or should be made from the individual tests? What about from the results in aggregate?
- What additional tests would you conduct if you had more data?

2.5 Option B: Bayesian Inference

- **Methods**
  - Draw a graphical model, clearly indicating which variables are observed. Provide descriptions of any hidden variables you're trying to estimate.
- **Results**
  - Summarize and interpret your results.
  - Estimate any uncertainty in your estimates, providing clear quantitative statements of the uncertainty in plain English.
- **Discussion**
  - Elaborate on the limitations of your methods.
  - If your inference procedure had trouble converging, can you explain why?
  - Did you try other formulations / graphical models? If so, what worked or didn't work about each one?
  - What additional data would be useful for answering this question, and why? How would you add it to the graphical model?

2.6 Option C: Prediction with GLMs and nonparametric methods

- **Methods**
  - Describe what you're trying to predict, and what features you're using. Justify your choices.
  - Describe the GLM you'll be using, justifying your choice. Describe any assumptions being made by your modeling choice.
  - Describe the nonparametric method you'll be using, justifying your choice. Describe any assumptions being made by your modeling choice.
  - How will you evaluate each model's performance?
- **Results**
  - Summarize and interpret your models.
  - Estimate any uncertainty in your GLM predictions, providing clear quantitative statements of the uncertainty in plain English.
- **Discussion**
  - Which model performed better, and why? How confident are you in applying this to future datasets?

  - Explain the results from both models. You may choose to not provide explanations, but you must justify this choice.
  - Elaborate on the limitations of your models.
  - What additional data would be useful for improving your models?

2.7 Option D: Causal Inference

- **Methods**
  - Describe which variables correspond to treatment and outcome.
  - Describe which variables (if any) are confounders. If the unconfoundedness assumption holds, make a convincing argument for why.
  - What methods will you use to adjust for confounders?
- **Results**
  - Summarize and interpret your results, providing a clear statement about causality (or a lack thereof) including any assumptions necessary.
  - Where possible, discuss the uncertainty in your estimate and/or the evidence against the hypotheses you are investigating.
- **Discussion**
  - Did your data exhibit Simpson's paradox?
  - Elaborate on the limitations of your methods.
  - What additional data would be useful for answering this causal question, and why?
  - How confident are you that there's a causal relationship between your chosen treatment and outcome? Why?

2.8 Conclusions

- Summarize your key findings
- Based on your results, suggest a call to action. What interventions, policies, real-world decisions, or action should be taken in light of your findings?
- Did you merge different data sources? What were the benefits and/or consequences of combining different sources?
- What limitations are there in the data that you could not account for in your analysis?
- What future studies could build on your work?

3 Project Deliverables

You must work in groups of four. Each group will submit one set of the following deliverables, submitted to Gradescope as a zip file.

3.1 Project Proposal (Due 4/30)

By April 30th at 11:59 PM PDT, please fill out the Project Proposal Form to indicate the dataset and research questions you'll be working on.

3.2 Written Report

You must submit a typed PDF document that contains each of the sections described in "Section Guidelines". Your report should be between 4-6 pages (single-spaced, 12 point font), excluding tables, figures, and references. All mathematical equations must be rendered properly in LaTeX, Equation Editor, or similar.

Your report should be a proper written document: you cannot just submit a printed Jupyter notebook (including data sources, code, outputs, etc.) We highly recommend using Overleaf or Google Docs. For your convenience, we have posted a list of LaTeX resources on Piazza.

If relevant, include a reference page with citations of all outside sources used.

All figures and tables should be appended to the end of your written report. Clearly label all figures and include informative captions. These labels should be used to reference figures and tables in your written report. Refer to this guide for instructions on inserting images into a LaTeX file.

3.3 Jupyter Notebook

You must submit a single notebook that contains all the code run for the project. Your code should be clear and well-documented. Please label each section of code in markdown.

Your results should be completely reproducible from the code you submit. For all random processes, we recommend that you set a seed to ensure that your results are consistent when your code is rerun. If you use a nonstandard library (in other words, any library that hasn't been used in the course so far), please include installation code.

3.4 Dataset (for external datasets)

If you choose to use your own data set, please submit the files. Make sure that the file names and files paths correspond to how you load the data in your Jupyter notebooks.

3.5 Team Member Assessment

When you submit your project, you will be required to fill out a form to quantify the contributions that each team member (including yourself) made to the project.

4 Datasets

We highly recommend selecting from the listed datasets; course staff has ensured that these datasets are sufficient in satisfying the project requirements. You may supplement the suggested data with secondary datasets (e.g., US Census, American Community Survey). If you bring in additional data, please reference the source in your write-up and submit the files in your ZIP file. (See Section 3.4).

If you choose to use an outside dataset, please refer to the suggested guidelines. Please note that staff will be less equipped to provide assistance with data outside the recommended list.

4.1 Dataset 1: Chronic Disease and Air Quality

Data Description:

The Center for Disease Control and Prevention (CDC) maintains the U.S. Chronic Disease Indicators dataset, containing state-specific data of chronic illness prevalence as well as relevant policies and regulations.

The CDC also publishes daily air quality data through the National Environmental Public Health Tracking Network to monitor environmental exposures.

This data is very large; we recommend subsetting the data to a particular geographic region and/or chronic illness to make the data easier to work with.

Potential Directions:

1. What has been the impact of substance regulation on chronic disease onset?
2. How do levels of particulate matter and ozone affect the onset of chronic illnesses, such as asthma?
3. Are there any geographical trends in air pollution and/or chronic illness?

Relevant datasets:

1. CDC: Annual State-Level U.S. Chronic Disease Indicators
   (a) Filtered for COPD
   (b) Filtered for Asthma
   (c) Filtered for Cardiovascular Disease
   (d) Filtered for Tobacco
2. CDC: Daily Census-Tract PM2.5 Concentrations
3. CDC: Daily Census-Tract Ozone Concentrations

Supplemental Readings:

1. Public health impact of global heating due to climate change: potential effects on chronic non-communicable diseases (2009)
2. Past Racist "Redlining" Practices Increased Climate Burden on Minority Neighborhoods (2020)

4.2 Dataset 2: Transportation, Mobility, and Infrastructure

Data Description:

The Bureau of Transportation Statistics (under the U.S. Department of Transportation) publishes monthly data transportation utilization and spending. The dataset includes information on airline traffic, transit ridership, transportation employment, construction spending, and transborder movement.

At the start of the COVID-19 pandemic, many tech companies released aggregated, anonymized data on global mobility patterns for researchers to study the effect of movement on disease transmission. Since early 2020, Google has maintained the Community Mobility Data to track differences in movement to/from workplaces, retail stores, and other community centers compared to pre-pandemic baseline levels.

Potential Directions:

1. How has infrastructure spending changed over time? What categories have seen substantial increases/decreases in investment in recent years?
2. Is there a relationship between investment and utilization?
3. How did human mobility change in response to the pandemic? How did those changes impact trends in transportation?

Relevant datasets:

1. Bureau of Transportation Statistics: Monthly Transportation Statistics
2. Google: Daily Community Mobility Data

Supplemental Readings:

1. Biden Details $2 Trillion Plan to Rebuild Infrastructure and Reshape the Economy (2021)
2. Impact of COVID-19 pandemic on mobility in ten countries and associated perceived risk for all transport modes (2021)

4.3 Dataset 3: Primary Election Endorsements and Financing

Data Description:

FiveThirtyEight compiled a dataset with information on numerous primary elections (primary elections determine the candidates that each political party nominates for the general election) in 2018. The data contains information about each candidate, including their political leanings, endorsements, and gender, race, and veteran identities. The dataset also reports outcomes for each election.

The Federal Election Commission (FEC) publishes campaign financing data, including the amount of data raised and spent by each candidate. Public data includes donor names, contribution amounts, and dates of contribution.

Potential Directions:

1. Is there a relationship between the number of unique donations a candidate received and the proportion of the vote they received in the primary?

2. What type of candidates did Joe Biden, Donald Trump, or Bernie Sanders endorse? What candidates were popular among different special interest groups?
3. What are the characteristics of the most contentious elections?

Relevant datasets:

1. FiveThirtyEight: 2018 Primary Candidate Endorsements
2. Federal Election Commission: 2018 Campaign Financing Data

Supplemental Readings:

1. The Persuasion Effects of Political Endorsements (2016)
2. How Money Affects Elections (2018)

4.4 Guidelines for External Data

If you choose to use an external dataset, the data must meet the following guidelines:

1. Data cannot contain sensitive and/or identifying data. If you do not have permission to share data with course staff, you are not allowed to use it for this project.
2. The data source must be known. Why was the data collected? Who conducted/funded data collection? When was the data measured or recorded? How was data collected?
3. At minimum, the data should include 2 numeric and 2 categorical variables.
4. Your dataset should have a sufficient number of observation. This is a fairly subjective measure: please check with staff if you are concerned that your data is too small.

4.5 External Data Sources

Listed below are some suggested data sources. This is not a comprehensive list, nor can we guarantee that data found on these sites will meet the guidelines for this project.

1. Humanitarian Data Exchange
2. Gapminder
3. World Bank
4. WHO
5. UNICEF
6. UC Irvine Machine Learning Repository
7. Google Data Repository
8. AWS Data Repository
9. FiveThirtyEight
10. Kaggle

5 Grading

The final project is worth 20% of your overall grade. Each section of the project will be weighted as follows:

1. Data Overview (12%)
2. EDA (10%)
3. Research Question 1 (30%)
   (a) 10% for methods
   (b) 10% for results
   (c) 15% for discussion
4. Research Question 2 (30%)
   (a) 10% for methods
   (b) 10% for results
   (c) 15% for discussion
5. Conclusion (12%)
6. Project Submission and Formatting (6%)
   (a) 2% for submitting project proposal
   (b) 2% for submitting group evaluations (additional adjustments may be made based on individual contributions)
   (c) 2% for following all other submission instructions

These values are subject to change. A more detailed rubric will be released on Piazza at a later date.